
Can Vision Models Read?

Probing Linguistic Structure via JEPA on Text Screenshots

Bariş Cem Bakay¹ Halil İbrahim Kanpak¹

Abstract

We ask whether vision models acquire linguistic structure when their training data consists of rendered text. As a probe, we use retrieval over a 2 000-page held-out slice of Wikipedia screenshots: given a 224×224 crop, identify its source page. We compare six encoders trained with five distinct objectives, including a joint-embedding predictive architecture (I-JEPA) that is the canonical image-only predictive baseline. The finding is sharp: image-text contrastive encoders (CLIP, SigLIP) retrieve at $R@10 = 0.78$ and 0.62 zero-shot, while image-only predictive (I-JEPA), self-supervised contrastive (DINOv2), supervised classification (ImageNet ViT), and random-init plain ViT all sit at $R@10 \leq 0.058$. I-JEPA in particular scores *below* random initialisation. Across the six encoders, mutual- k NN alignment with text encoders rank-correlates with retrieval $R@10$ at Spearman $\rho = 0.83$ ($p = 0.04$), supporting the Platonic Representation Hypothesis as the explanatory axis. Fine-tuning DINOv2 with a SALAD aggregator on 50 k pages reaches $R@10 = 0.915$, exceeding zero-shot CLIP, but every fine-tuned checkpoint exhibits a layout-fingerprint failure mode: erasing the top 15–30 % of every test image drops in-distribution Protocol-A $R@10$ by up to 0.29 while *improving* cross-set Phase-2 $R@1$ by up to $6.7\times$ on the same checkpoints. Zero-shot CLIP and SigLIP move in the opposite direction ($R@10$ *rises* by 0.07–0.11 under blanking), falsifying the title-OCR explanation for their zero-shot edge. A linear adapter trained on (image, BERT-of-title) pairs lifts I-JEPA from $R@10 = 0.015$ to 0.197, recovering latent text alignment that is invisible without the projection.

1. Introduction

Documents combine glyph sequences with visual layout. A vision encoder trained on rendered text could in principle acquire two distinct kinds of structure: (i) glyph-level reading (recognising words and characters as visual patterns), and (ii) layout-level structure (titles, infoboxes, lists, figures, captions). The two pipelines in Figure 1 operationalise the alternatives: the text pipeline OCRs the page and encodes characters; the visual pipeline encodes the screenshot directly. The visual pipeline is attractive when text extraction is unreliable (degraded scans, formula-heavy or multilingual content) or when the query is itself a visual crop with no text available, but it places the burden of “reading” on the encoder.

We probe whether modern vision encoders read in this sense by measuring how well they retrieve Wikipedia source pages from 224×224 crops. The probe is informative because the visual content *is* the text: an encoder that genuinely reads should match an encoder that operates on the underlying string. We compare five families:

- **Image-text contrastive** (CLIP, SigLIP) — learns by matching images and captions; sees text indirectly.
- **Image-only predictive** (I-JEPA) — the canonical joint-embedding predictive architecture; trains an image encoder to predict the latent representations of masked image blocks (Assran et al., 2023). No text supervision.
- **Image-only contrastive** (DINOv2) — view-contrastive distillation on natural images. No text supervision.
- **Supervised classification** (ImageNet ViT) — ViT-B/16 trained on ImageNet-21k+1k labels.
- **Random initialisation** — a frozen ViT-B/16 sanity floor.

Three questions:

Q1 (zero-shot reading). Which pretraining produces image representations from which the source page can be retrieved zero-shot?

¹Department of Computer Engineering, Koç University, Istanbul, Turkey. Correspondence to: Halil İbrahim Kanpak <hkanpak21@ku.edu.tr>.




Figure 1. The two retrieval paradigms compared in this paper. The *text* pipeline (top) extracts characters with an OCR tool and encodes them with a text encoder. The *visual* pipeline (bottom) feeds the rendered document image directly into a vision encoder. We study the visual pipeline because it bypasses OCR errors, multilingual tokeniser failures, and layout discarding, and because it admits image queries (e.g. a crop of an unknown page) for which no text exists.

Q2 (JEPA specifically). Does I-JEPA, an image-only predictive objective with no text supervision, acquire any linguistic structure detectable through retrieval?

Q3 (the Platonic axis). Does the Platonic Representation Hypothesis (Huh et al., 2024) — which predicts that encoders trained on diverse objectives converge to a shared abstract representation — explain the retrieval ordering quantitatively? We measure each image encoder’s mutual- k NN alignment with text encoders on the same content and ask whether alignment predicts retrieval performance.

Contributions. (1) A leak-free crop→page protocol with leave-one-out gallery aggregation. (2) A six-encoder zero-shot grid showing image–text contrastive encoders dominating the four image-only baselines, with I-JEPA sitting *below* random initialisation. (3) A Platonic alignment grid producing the rank correlation $\rho = 0.83$ between alignment and retrieval, statistically significant at $n = 6$. (4) A reproducible failure mode in every fine-tuned checkpoint: erasing the top 15–30 % of every test image *improves* cross-set transfer, indicating the model overweights a layout confound rather than reading content.

2. Related Work

Pixel-based language models and document AI. PIXEL (Rust et al., 2023), CLIPPO (Tschannen et al., 2023) and Pix2Struct (Lee et al., 2023) all train transformers on rendered text and demonstrate that pixel inputs can match or exceed tokenised text on downstream tasks. ColPali (Faysse et al., 2025) produces multi-vector page embeddings via late interaction and shows that direct visual retrieval over screenshots competes with text+OCR pipelines on the ViDoRe benchmark. We adopt a simpler single-vector setup to keep our six encoders comparable on identical infrastructure.

Self-supervised vision and post-hoc text alignment. DI-

NOV2 (Oquab et al., 2024) learns image features via view-contrastive distillation; I-JEPA (Assran et al., 2023) learns them by predicting masked-block representations. Jose et al. (Jose et al., 2024) showed that a small text encoder aligned post-hoc to a frozen DINOv2 reaches CLIP-level zero-shot classification, suggesting that self-supervised image features already contain a linguistically-decodable structure. We test the analogous claim for I-JEPA on document screenshots.

Image–text contrastive pretraining and OCR shortcuts. CLIP (Radford et al., 2021) and SigLIP (Zhai et al., 2023) are trained contrastively on image–caption pairs, with CLIP using softmax InfoNCE and SigLIP using a sigmoid pairwise loss. Both have been shown to read text rendered in images (Materzyńska et al., 2022; Lin et al., 2024); we exploit this in a title-blanking experiment.

Platonic alignment and probing. Huh et al. (Huh et al., 2024) predict cross-encoder alignment grows with model and data scale; Maniparambil et al. (Maniparambil et al., 2024) and Moayeri et al. (Moayeri et al., 2023) measure such alignment for natural-image vision encoders. We extend this measurement to a document-screenshot domain and use mutual- k NN, debiased CKA (Kornblith et al., 2019; Murphy et al., 2024), and Procrustes distance on identical 2 000-page eval slices.

3. Method

3.1. Data and crops

We use Tevatron/wiki-ss-corpus: ~ 1.2 M Wikipedia page screenshots paired with the page title and rendered body text. We hold a local cache of ~ 376 k pages. A non-overlapping cropper tiles each page into 490×490 px crops resized to 224×224 px. At this resolution body text is sub-pixel and illegible to a ViT patch encoder; only headings and the overall layout structure




Figure 2. Training. Each step samples four pages $\times$ four crops; same-page crops are positives, cross-page crops are negatives. The image encoder produces an L2-normalised descriptor; training minimises a multi-similarity contrastive loss (Wang et al., 2019). We unfreeze the last four ViT blocks (of twelve) and train a head on top.

remain visually recoverable. “Reading” in this paper therefore refers to heading-level text and visual layout, not fluent body content.

3.2. Training (fine-tunes)

We train each (encoder, head) configuration with the procedure of Figure 2. Each batch contains 4 pages $\times$ 4 crops = 16 images. Within a batch, crops sharing a page are positives; otherwise negatives. The loss is multi-similarity (Wang et al., 2019). The optimiser is AdamW with weight decay 10^{-4} , gradient clip 1.0, linear warmup over 5 % of steps followed by cosine decay. Backbone learning rate is 10^{-5} for DINOv2 and 10^{-7} for CLIP and SigLIP (a higher rate causes catastrophic representation collapse on the latter two), head learning rate 5×10^{-4} throughout. Heads are either an MLP ($768 \rightarrow 512 \rightarrow 256$) or the SALAD optimal-transport aggregator (Izquierdo & Civera, 2024) producing an 8448-d descriptor. Only the head and the last four ViT blocks are trainable. Fine-tunes use 30 000 training pages $\times$ 3 epochs (5 625 optimiser steps), one seed; the saturation diagnostic uses 50 000 pages.




Figure 3. Anchor-pool benchmark (Phase 2). Each anchor crop is encoded with the vision encoder and ranked against a small pool (~ 7 images) of positives (other crops of the same page) and hard negatives (crops of unrelated pages).

3.3. Evaluation

We evaluate on a held-out slice of $P = 2000$ pages disjoint from training, with two complementary protocols.

Protocol A — crop $\rightarrow$ page (headline metric). For each query crop c on page p with crops $\{c_1, \dots, c_{n_p}\}$, the gallery vector for any page $q \neq p$ is the L2-normalised mean of all q ’s crop embeddings. The gallery vector for page p is recomputed leaving the query out: $g_p^{\setminus c} = \text{normalise}(\sum_i \phi(c_i) - \phi(c))$. We cosine-rank the P gallery vectors against $\phi(c)$ and report $R@k$ as the fraction of queries whose source page lands in the top k . Random $R@1 = 1/P = 5 \times 10^{-4}$. The leave-one-out aggregation forbids the trivial “find another crop of the same column” shortcut implicit in earlier crop-vs-crop formulations of this task.

Protocol B — anchor-pool transfer. The anchor-pool benchmark of Figure 3 uses 118 valid validation triplets from Wikipedia_SS.withanchors. For each anchor crop, the pool is the union of its positives and hard negatives (~ 7 images). We rank the pool by cosine similarity and report $R@1$. This measures cross-set transfer: anchors and pool images come from a different distribution of crops than what the model trained on.

3.4. Platonic alignment

For each (image encoder E_v , text encoder E_t) pair we encode the same 2 000-page eval slice with both halves: image features $\mathbf{X} \in \mathbb{R}^{N \times d_v}$ from page screenshots, text features $\mathbf{Y} \in \mathbb{R}^{N \times d_t}$ from page titles. We report mutual- k NN at $k = 10$ (the fraction of points sharing at least one of their top-10 neighbours under the two embeddings), debiased CKA (Kornblith et al., 2019; Murphy et al., 2024), and

Table 1. Protocol-A zero-shot retrieval at $P = 2000$. Random $R@1 = 5 \times 10^{-4}$. Encoder rows are coloured by pretraining family: image-text contrastive image-only self-supervised

Encoder	R@1	R@5	R@10	R@20	gap
CLIP ViT-B/16	0.548	0.715	0.779	0.833	+0.205
SigLIP ViT-B/16	0.348	0.541	0.624	0.703	+0.110
DINOv2 ViT-B/14	0.022	0.043	0.058	0.075	+0.010
I-JEPA ViT-H/14	0.002	0.008	0.015	0.021	−0.002
I-JEPA Text-Target	0.007	0.019	0.029	0.046	+0.004
ImageNet ViT-B/16	0.009	0.022	0.031	0.043	−0.002
Plain ViT-B/16	0.007	0.020	0.029	0.042	+0.000

Procrustes distance. The Platonic Hypothesis predicts a positive rank correlation between E_v ’s alignment and E_v ’s downstream cross-modal retrieval performance. We test this directly in §4.

4. Results

4.1. Zero-shot retrieval

Table 1 reports zero-shot Protocol-A. CLIP dominates at $R@10 = 0.779$, SigLIP follows at 0.624, and the four image-only encoders cluster near random ($R@10 = 0.015$ –0.058). Three observations follow:

- (i) The largest gap is between the image-text contrastive and the image-only families ($\Delta R@10 \approx 0.57$), not between encoders within a family. The supervised ImageNet ViT scores 0.031 versus 0.029 for the random-init plain ViT, so ImageNet classification supervision adds essentially nothing for this task.
- (ii) I-JEPA scores $R@10 = 0.015$, *below* the random-init plain ViT (0.029). Despite ViT-H/14 being roughly $5\times$ larger than the ViT-B encoders, image-only predictive pre-training out-of-the-box does not transfer to visual document retrieval. Q2’s optimistic reading is therefore falsified at the zero-shot level: a post-hoc text adapter (§5) is the only remaining path to rescue I-JEPA.
- (iii) The CLIP–SigLIP gap (0.155) is real but small relative to the CLIP–DINOv2 gap (0.721). The dominant predictor is whether the encoder saw paired image-text data, not the specific contrastive form (softmax vs. sigmoid).

4.2. Fine-tuning

Table 2 reports the fine-tuned grid. DINOv2 + SALAD on 30 k pages reaches $R@10 = 0.880$ and on 50 k pages reaches 0.915, exceeding zero-shot CLIP (0.779). The MLP head on the same DINOv2 backbone scores 0.787, so the SALAD aggregator contributes +9.3 pp $R@10$ over

Table 2. Protocol-A retrieval after last-4-block fine-tuning, single seed. “ ΔZS ” is $R@10$ minus the encoder’s zero-shot baseline. Yellow rows are fine-tuned variants of the encoder above them.

Encoder + head	R@1	R@10	ΔZS
DINOv2 + MLP (30k)	0.351	0.787	+0.729
DINOv2 + SALAD (30k)	0.512	0.880	+0.822
DINOv2 + SALAD (50k)	0.569	0.915	+0.857
CLIP + SALAD (30k)	0.601	0.724	−0.055
CLIP + MLP (30k) [†]	0.000	0.000	collapsed

[†] Embeddings collapsed to a constant.

an MLP at the matched backbone — consistent with the prior place-recognition result of (Izquierdo & Civera, 2024).

CLIP fine-tuning at the conservative $\eta_{bb} = 10^{-7}$ hurts more than it helps. CLIP+SALAD loses 5.5 pp $R@10$ from its own zero-shot baseline; the same-page minus diff-page sanity gap shrinks from +0.205 to +0.032, indicating the fine-tune sharpened the top-1 match ($R@1 +5.3$ pp) but compressed the overall similarity geometry. CLIP+MLP collapsed to a constant output by step 400. Training CLIP at the LR that worked for DINOv2 (10^{-5}) caused immediate divergence in earlier attempts; 10^{-7} was chosen to avoid that, but evidently the CLIP+MLP combination cannot navigate the resulting flat-loss basin. A multi-LR sweep is the proper next step but is deferred.

4.3. Native Legible Protocol (224 → 224)

To analyze the models’ capability of reading legible, native-resolution text (rather than downsampled headlines), we establish the Native Legible Protocol (224 → 224). Here, Wikipedia screenshots are cropped without downsampling, keeping body text sharp and legible. Table 3 reports the results for the main baselines, while Table 4 isolates different I-JEPA pretraining configurations. While the family ranking (contrastive vs. SSL) is preserved, we observe that zero-shot retrieval improves for all models since text templates are not distorted. The MAE baseline collapses under standard regression but recovers under InfoNCE contrastive training to yield 0.098 $R@10$, which further improves to 0.143 when fine-tuned with random title-region masking.

4.4. Vision-to-vision adapter capacity

To isolate the effect of adapter capacity from backbone fine-tuning, we sweep linear, MLP, and bottleneck heads over frozen backbones (Table 5). We observe that CLIP completely collapses to a constant representation ($R@10 \approx 0.003$) regardless of adapter capacity, stemming from its lack of fine-grained patch variance on document crops. Conversely, DINOv2 and I-JEPA learn meaningful projections, with the simplest linear adapter yielding the best performance (0.229 for DINOv2), suggesting that higher-capacity

Table 3. Protocol-A retrieval R@10 under the Native Legible Protocol (224 → 224, no downsampling) for different pretrainings. Under this protocol, body text is fully legible to the models. Missing entries indicate configurations not yet trained at this resolution.

Backbone / Pretraining	Zero-Shot	MLP	SALAD
Perfect-Text Upper Bounds (Linguistic Ceiling)			
BERT (mean-pool)	0.938	—	—
BERT (CLS)	0.747	—	—
Visual Models			
CLIP ViT-B/16	0.736	—	—
SigLIP ViT-B/16	0.697	—	—
DINOv2 ViT-B/14 (CLS)	0.116	—	—
ImageNet ViT-B/16	0.084	—	—
I-JEPA ViT-H/14	0.054	—	—
Plain ViT-B/16	0.042	—	—
MAE ViT-B/16	0.036	0.098 [†]	—
MAE ViT-B/16 (title-masked)	—	0.143 [†]	—

[†] Fine-tuned using InfoNCE contrastive training to BERT-mean.

Table 4. Protocol-A retrieval R@10 under the Native Legible Protocol (490 → 490, no downsampling) specifically comparing different I-JEPA pretraining configurations. All variants are pre-trained on 90k Wikipedia screenshots at native 490×490 resolution with all 32 ViT-H/14 blocks unfrozen. Head fine-tunes use 30k pages. Entries marked TBD are from runs currently in progress.

I-JEPA Variant	Zero-Shot	MLP	SALAD
Baseline (Image-Only)	0.029	TBD	TBD
Text-Target (Body) [†]	0.029	TBD	TBD

[†] Backbone predicts frozen BERT representations of the page body text. A title-targeted variant is deferred to future work.

heads overfit the frozen visual features.

4.5. Platonic alignment predicts retrieval

Table 6 reports image-encoder × text-encoder mutual- k NN@10 alignment. CLIP and SigLIP align with text encoders at 0.09–0.14; the four image-only encoders align at 0.007–0.027. Figure 4 plots, for each image encoder, its column-maximum text alignment against its zero-shot Protocol-A R@10. The Spearman correlation is 0.83 ($p = 0.04$, $n = 6$): statistically significant even at this small sample size. The two image-text contrastive encoders sit in the upper-right corner; the four image-only encoders sit in the lower-left.

This is the central empirical claim of the paper. The kind of pretraining the encoder received is a strong predictor of its retrieval performance, mediated by its alignment with text representations on the same content. The Platonic Hypothesis is supported in this domain: encoders close to the text-encoder geometry retrieve text-heavy visual documents better.

Table 5. Protocol-A Phase 1 R@10 for the frozen-backbone vision-to-vision adapter grid. Backbones are frozen, and only the adapter head is trained for 3 epochs. CLIP fails completely due to lack of local variance, whereas DINOv2 and I-JEPA perform best with a simple linear adapter.

Backbone	Linear	MLP	Bottleneck
DINOv2	0.229	0.207	0.184
I-JEPA	0.180	0.157	0.146
CLIP	0.003	0.003	0.003

Table 6. Mutual- k NN@10 between image and text encoders on the same 2 000-page eval slice. Bold = highest entry per row. Row colours match Table 1.

image \ text	BERT	MiniLM	CLIP-T	SigLIP-T
CLIP	0.094	0.120	0.138	0.132
SigLIP	0.091	0.108	0.123	0.131
DINOv2	0.027	0.021	0.019	0.026
I-JEPA	0.019	0.014	0.016	0.018
ImageNet ViT	0.018	0.013	0.013	0.019
Plain ViT (rand)	0.008	0.007	0.007	0.009

4.6. Title-region ablation

We re-encode the same 2 000-page eval slice with the top 15 % (and 30 %) painted white. Three signs emerge. Zero-shot CLIP and SigLIP *improve* under blanking (+0.073 and +0.106 R@10); image-only encoders move by ≤ 0.003 . This falsifies the H-OCR explanation for CLIP’s zero-shot edge — the title acts as a low-information template common to every page, and removing it yields cleaner content-driven matches. Fine-tuned DINOv2 *loses* heavily (−0.16 to −0.29 R@10); the 50 k checkpoint falls from 0.915 to 0.673. CLIP+SALAD moves only −0.037. Fine-tuning has taught DINOv2 to bind page identity to the title-bearing top region. Phase-2 reverses the sign on the same checkpoints (Table 8): the title-region features that help in-distribution Protocol-A are the same confound that hurts cross-set transfer, with up to $6.7\times$ R@1 improvement under 30 % blanking. The encoder is not unable to read the body region; it is that the title is so reliably informative on Wikipedia that the model bypasses the body. Random title-region masking during training is the obvious mitigation.

4.7. Text adapter capacity sweep

Q2’s zero-shot verdict (Table 1) is sharp: I-JEPA does not produce text-retrievable image representations out of the box. We test a softer version of Q2 — do visual features encode latent text-aligned structure recoverable by post-hoc adapters of varying capacity? We sweep across linear, low rank, bottleneck, and MLP adapters mapping frozen I-JEPA (1280-d) and CLIP (768-d) features into BERT space (Table 9).




Figure 4. Each point is one image encoder. The x -axis is its maximum mutual- k NN@10 against the four text encoders (column max of Table 6); the y -axis is its zero-shot Protocol-A R@10 (Table 1). Marker shape and pastel colour code the pretraining family. Spearman $\rho = 0.83$ ($p = 0.04$, $n = 6$).

Table 7. Protocol-A R@10 with the top 15 % of every page painted white before cropping (b-15) versus the original eval (orig.). Top block: zero-shot encoders. Bottom block: fine-tuned checkpoints. Sign of Δ is the central observation: zero-shot CLIP/SigLIP *gain* from blanking (negative confound removed); trained DINOv2 *loses* substantially (the model learned to depend on the title region).

Encoder	orig. R@10	b-15 R@10	Δ
CLIP zero-shot	0.779	0.852	+0.073
SigLIP zero-shot	0.624	0.731	+0.106
DINOv2 zero-shot	0.058	0.058	−0.000
I-JEPA zero-shot	0.015	0.012	−0.003
ImageNet zero-shot	0.031	0.032	+0.001
Plain-ViT zero-shot	0.029	0.030	+0.001
DINOv2-MLP-30k	0.787	0.494	−0.292
DINOv2-SALAD-30k	0.880	0.720	−0.161
DINOv2-SALAD-50k	0.915	0.673	−0.242
CLIP-SALAD-30k	0.724	0.687	−0.037

While CLIP completely failed at the pure vision-to-vision contrastive task due to collapsed variance on white document backgrounds, it excels at predicting BERT text embeddings, jumping to ~ 0.717 R@10. This is because CLIP’s visual backbone was explicitly pre-trained to align with text encoders. The simple linear or low-rank adapters are sufficient for CLIP to reach peak performance. In contrast, I-JEPA peaks at ~ 0.219 R@10 and requires the highest capacity deep MLP adapter to do so. Because I-JEPA was trained via masked image modeling, its features require heavy non-linear transformations to align with text semantics.

Table 8. Phase-2 cross-set R@1 with the top fraction of every anchor and pool image painted white. The same fine-tuned checkpoints that lose Protocol-A R@10 under blanking (Table 7) *gain* Phase-2 R@1 by the same information removal. CLIP+MLP-30k embeddings collapsed; all predictions are identical.

Run	orig. R@1	b-15 R@1	b-30 R@1
DINOv2-MLP-30k	0.169	0.712	0.763
DINOv2-SALAD-30k	0.212	0.686	0.737
DINOv2-SALAD-50k	0.110	0.661	0.737
CLIP-SALAD-30k	0.551	0.686	0.720
CLIP-MLP-30k [†]	0.932	0.932	0.932

[†] Collapsed model.

Table 9. Image-to-text retrieval Phase 1 R@10 on a held-out 1 000-page set. Image embeddings from frozen backbones are mapped into BERT[CLS] space (768-d) representing the page title. Adapters are trained on 4 000 screenshot-title pairs. CLIP dominates I-JEPA, as its pretraining inherently aligns visual features to text, whereas I-JEPA struggles to map pure visual patch features to text semantics.

Adapter Architecture	I-JEPA	CLIP
Baseline (no adapter)	0.009	0.008
Linear	0.214	0.715
Low Rank ($r = 16$)	0.172	0.471
Low Rank ($r = 64$)	0.212	0.661
Low Rank ($r = 256$)	0.213	0.717
Bottleneck ($r = 64$)	0.202	0.666
Bottleneck ($r = 256$)	0.207	0.712
MLP ($h = 2048$)	0.213	0.709
Deep MLP ($h = 2048$)	0.219	0.637

4.8. Late interaction is not the bottleneck-breaker

Does single-vector pooling discard retrievable structure that a patch-level scorer would recover? Table 10 reports late interaction on frozen CLIP and DINOv2 at $P = 500$. Late-interaction CLIP gains only +0.04 R@10 over its single-vector baseline and loses R@1 from 0.705 to 0.178: the sum-of-max scoring bunches plausible pages at similar totals, so the source page lands in the top-5 but rarely first ($R@5=R@10=R@20=0.928$ means failures miss entirely rather than rank 6–20). DINOv2 late-interaction stays at chance (within 0.01 of single-vector DINOv2 at the same pool). Patch-level scoring does not rescue an encoder that lacks text-aligned features; the bottleneck is the encoder, not the aggregator.

4.9. LEACE causal title-erasure

Pixel-blanking is correlational. We isolate the *title-direction* in feature space and erase it surgically: the rank-1 PCA projection of the per-crop deltas serves as a 1-d protected attribute for the LEACE eraser. Table 11 reports the result.

For both image-text contrastive encoders, LEACE-erasing

Table 10. Mean-pool single-vector vs. ColPali (Faysse et al., 2025) late-interaction at matched $P = 500$. Late interaction keeps all 196 (CLIP) or 256 (DINOv2) patch tokens per crop and scores by summed-max-cosine.

Encoder + scoring	R@1	R@5	R@10	R@20
CLIP single-vector	0.705	0.843	0.890	0.924
CLIP late-interaction	0.178	0.928	0.928	0.928
DINOv2 late-interaction	0.023	0.052	0.066	0.085

Table 11. Closed-form LEACE (Belrose et al., 2023) erasure of the rank-1 PCA projection of per-crop deltas $\phi(\text{orig}) - \phi(\text{blank})$, $P=2000$. Higher-rank Z over-erases (rank(Z)= d collapses the feature space). Last column: explained variance of the rank-1 PCA.

Encoder	R@10 orig	R@10 blank	R@10 LEACE	Expl. var.
CLIP zero-shot	0.779	0.823	0.865	0.865
SigLIP zero-shot	0.611	0.716	0.775	0.775
DINOv2 zero-shot	0.058	0.056	0.068	0.068
DINOv2-SALAD-50k trained	0.911	0.670	0.912	0.912

a single linear axis raises retrieval *more* than the pixel intervention: CLIP +0.084 vs +0.044, SigLIP +0.164 vs +0.105. The title-direction’s dominant axis carries enough of the encoder’s confounding signal that surgical removal beats the physical one. Zero-shot DINOv2 is unmoved by either: the title direction barely exists as a structured axis in its feature space.

Trained DINOv2-SALAD-50k splits the analysis. Pixel-blanking drops its R@10 from 0.911 to 0.670 (−0.241), but rank-1 LEACE leaves retrieval untouched (0.912). A rank sweep {4, 16, 64, 256, 1024} shows the title-region binding lives in a 16–64-dimensional subspace: rank-4 LEACE drops R@10 by 0.030, rank-16 by 0.136, rank-64 by 0.266 (matching pixel-blanking), and rank-256 by 0.385. Pretraining gives CLIP and SigLIP a clean low-rank title confound a single axis can erase; fine-tuning on a single visual corpus spreads the analogous binding across an order-of-magnitude wider subspace of the 8 448-d SALAD descriptor.

5. Discussion and Limitations

All training uses a single random seed. The cross-family ordering is large enough to read past seed noise (closest gap $\Delta R@10 = 0.093$, well outside typical jitter), but within-family comparisons (CLIP-SALAD vs. CLIP-MLP) need multi-seed reruns. Pool size $P=2000$ is small versus industrial 10^5 – 10^6 ; numbers are within-table comparable, not cross-paper. At 224^2 input, body text is sub-pixel, so “reading” here means heading-level text plus layout, not fluent content. The I-JEPA adapter R@10 = 0.197 is evidence of latent alignment, not a practical recipe; matching CLIP-level retrieval would need a non-linear adapter and larger

paired-data training.

6. Conclusion

We compared six image encoders trained with five distinct pretraining objectives on a leak-free crop→page retrieval task over Wikipedia screenshots. Image–text contrastive pre-training (CLIP, SigLIP) dominates zero-shot retrieval over image-only self-supervision (DINOv2, I-JEPA), supervised classification (ImageNet ViT), and random initialisation. Across the six encoders, alignment with text encoders rank-correlates with retrieval performance at Spearman $\rho = 0.83$, supporting the Platonic Representation Hypothesis on this domain. Fine-tuning DINOv2 with the SALAD aggregator on 50 k pages reaches R@10 = 0.915, exceeding zero-shot CLIP, but introduces a universal pathology: the fine-tuned model overweighs the title region, dropping in-distribution SigLIP A R@10 by up to 0.29 under blanking while improving cross-set Phase-2 R@1 by up to 6.7×. Zero-shot CLIP and SigLIP show the opposite sign under the same blanking, falsifying the simple title-OCR explanation for their zero-shot edge. A 1280×768 linear adapter trained on (image, BERT-title) pairs lifts I-JEPA from near-chance to R@10 = 0.197, indicating image-only predictive features encode latent text-aligned structure that is not zero-shot accessible but is recoverable by a low-capacity linear projection.

References

- Assran, M., Duval, Q., Misra, I., Bojanowski, P., Vincent, P., Rabbat, M., LeCun, Y., and Ballas, N. Self-supervised learning from images with a joint-embedding predictive architecture. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023.
- Belrose, N., Schneider-Joseph, D., Ravfogel, S., Cotterell, R., Raff, E., and Biderman, S. LEACE: Perfect linear concept erasure in closed form. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Faysse, M., Sibille, H., Wu, T., Omrani, B., Viaud, G., Hudelot, C., and Colombo, P. ColPali: Efficient document retrieval with vision language models. In *International Conference on Learning Representations*, 2025. arXiv:2407.01449.
- Huh, M., Cheung, B., Wang, T., and Isola, P. The platonic representation hypothesis. In *International Conference on Machine Learning*, 2024. arXiv:2405.07987.
- Izquierdo, S. and Civera, J. Optimal transport aggregation for visual place recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 17658–17668, 2024. arXiv:2311.15937.

- Jose, C. et al. DINOv2 meets text: A unified framework for image- and pixel-level vision-language alignment. Meta AI Technical Report, 2024.
- Kornblith, S., Norouzi, M., Lee, H., and Hinton, G. Similarity of neural network representations revisited. In *International Conference on Machine Learning*, 2019.
- Lee, K., Joshi, M., Turc, I., Hu, H., Liu, F., Eisenschlos, J. M., Khanelwal, U., Shaw, P., Chang, M.-W., and Toutanova, K. Pix2struct: Screenshot parsing as pretraining for visual language understanding. In *International Conference on Machine Learning*, 2023.
- Lin, Y., Wang, C., Wu, H., et al. Parrot captions teach CLIP to spot text. In *European Conference on Computer Vision (ECCV)*, 2024.
- Maniparambil, M. et al. Do vision and language encoders represent the world similarly? In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.
- Materzyńska, J., Torralba, A., and Bau, D. Disentangling visual and written concepts in CLIP. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (Oral)*, 2022.
- Moayeri, M. et al. Text2concept: Concept activation vectors directly from text. In *CVPR XAI4CV Workshop*, 2023.
- Murphy, A. et al. Why is the CKA similarity index misleading and how to fix it. In *ICLR Re-Align Workshop*, 2024.
- Oquab, M., Darcet, T., Moutakanni, T., Vo, H. V., Szafraniec, M., Khalidov, V., Fernandez, P., Haziza, D., Massa, F., El-Nouby, A., et al. DINOv2: Learning robust visual features without supervision. *Transactions on Machine Learning Research*, 2024.
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., and Sutskever, I. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning*, 2021.
- Rust, P., Lotz, J. F., Bugliarello, E., Salesky, E., de Lhoneux, M., and Elliott, D. Language modelling with pixels. In *International Conference on Learning Representations*, 2023. arXiv:2207.06991.
- Tschannen, M., Mustafa, B., and Houlsby, N. CLIPPO: Image-and-language understanding from pixels only. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023.
- Wang, X., Han, X., Huang, W., Dong, D., and Scott, M. R. Multi-similarity loss with general pair weighting for deep metric learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019.
- Zhai, X., Mustafa, B., Kolesnikov, A., and Beyer, L. Sigmoid loss for language image pre-training. In *International Conference on Computer Vision (ICCV)*, 2023.